

Inference runtime vs. gate layer placement
(MrBERT-30%, SNLI, A100, hard deletion)

